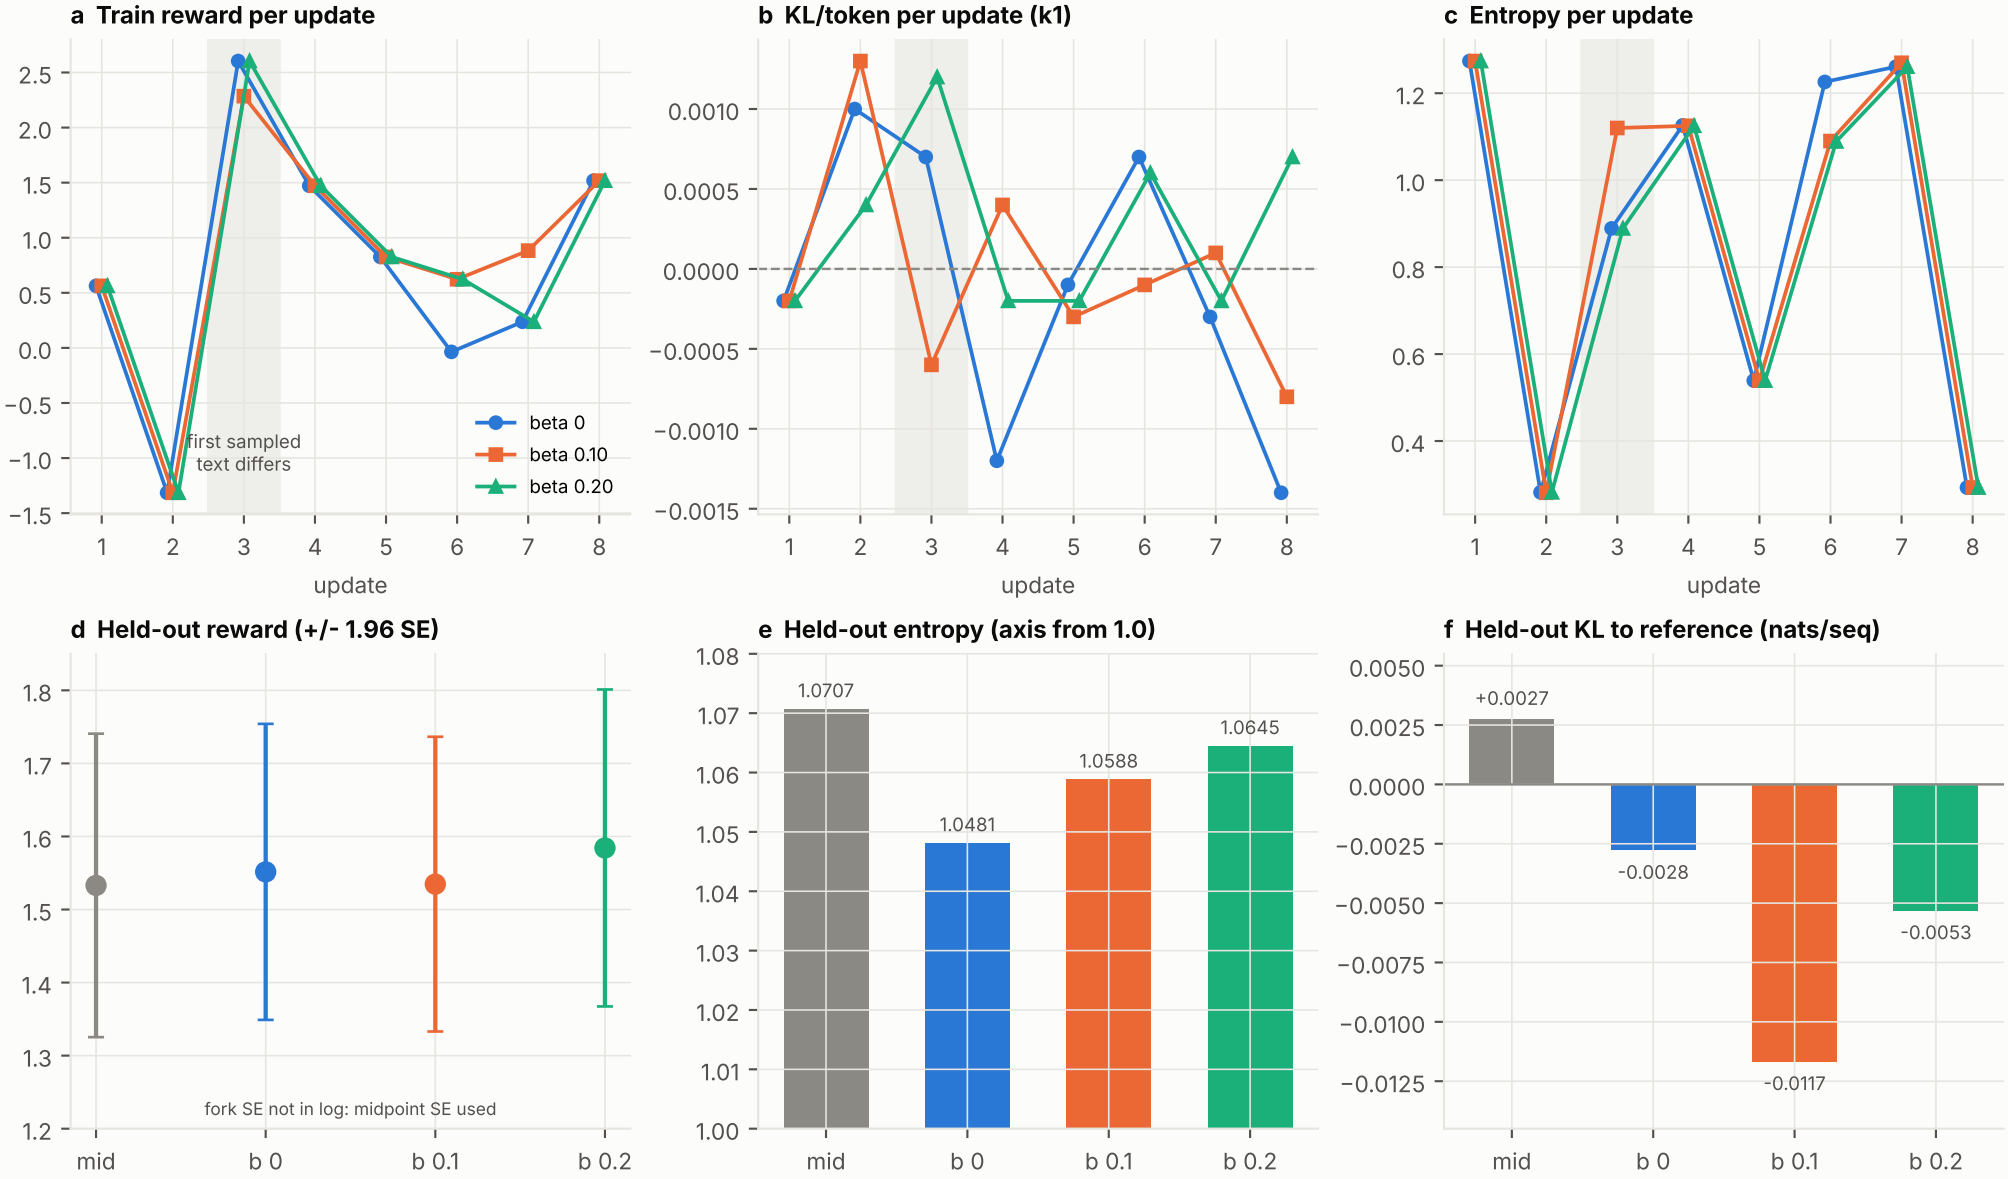
a Train reward per update
2.5
2.0
1.5
1.0
0.5
0.0
−0.5
−1.0
−1.5
first sampled
text differs
beta 0
beta 0.10
beta 0.20
1
2
3
4
5
6
7
8
update
b KL/token per update (k1)
0.0010
0.0005
0.0000
−0.0005
−0.0010
−0.0015
update
c Entropy per update
1.2
1.0
0.8
0.6
0.4
update
d Held-out reward (+/- 1.96 SE)
1.8
1.7
1.6
1.5
1.4
1.3
1.2
fork SE not in log: midpoint SE used
mid
b 0
b 0.1
b 0.2
e Held-out entropy (axis from 1.0)
1.08
1.07
1.06
1.05
1.04
1.03
1.02
1.01
1.00
1.0707
1.0481
1.0588
1.0645
mid
b 0
b 0.1
b 0.2
f Held-out KL to reference (nats/seq)
0.0050
0.0025
0.0000
−0.0025
−0.0050
−0.0075
−0.0100
−0.0125
+0.0027
-0.0028
-0.0117
-0.0053
mid
b 0
b 0.1
b 0.2