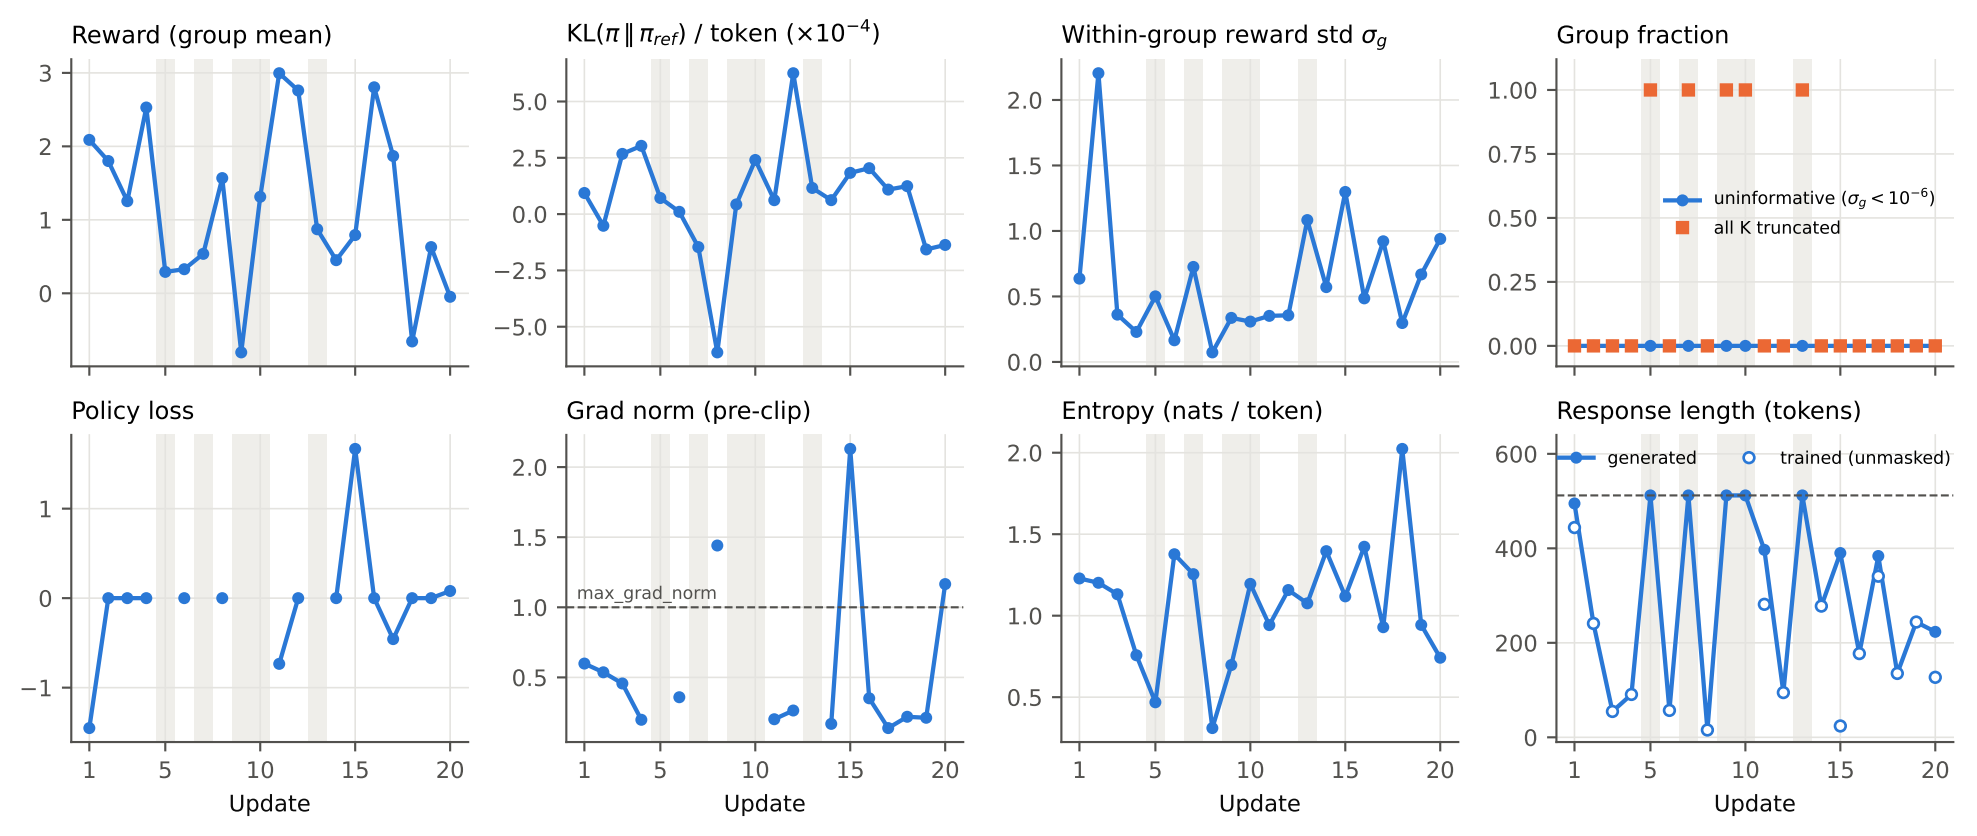
Reward (group mean)
KL($\pi \| \pi_{ref}$) / token ($\times 10^{-4}$)
Within-group reward std σ_g
Group fraction
uninformative ($\sigma_g < 10^{-6}$)
all K truncated
Policy loss
Grad norm (pre-clip)
max_grad_norm
Entropy (nats / token)
Response length (tokens)
generated
trained (unmasked)
Update